Against the Grey Template

A founder's briefing on AI, design, and where the business actually is

---

The opportunity, in short

Generation is now free. Any of a dozen tools will turn a sentence into a working screen in seconds. That's exactly the problem, and it's also the opening. When creation is free, the thing that's still scarce — and still worth paying for — is **judgment**: the ability to look at ten competent options and refuse nine of them because they're wrong for this brand. "Taste is defined by what you refuse to do." A model, by construction, doesn't refuse anything; it returns the statistical average of everything it's seen. That gap is where a defensible business lives.

Three things make this the right moment. Generation got good and got everywhere — Bolt hit ~$40M ARR in five months, Lovable went $100M→$200M ARR in about a year at a $6.6B valuation, and Vercel's v0, Figma's own agent, and Google's Stitch all ship credible output in seconds. The ceiling on machine *judgment* is now measurable rather than just felt — the best vision models agree with professional designers on "which is better" barely more than 70% of the time, and are close to useless at pointing out exactly where a design fails. And the infrastructure to encode real craft into a system — perceptually-correct color math (OKLCH), a vendor-neutral design-token standard (W3C's spec went stable in October 2025) — finally exists. Adoption is still early: McKinsey found only ~23% of organizations have scaled any agentic AI system past a pilot. The tooling, the ceiling, and the whitespace all lined up at once.

Slop is a P&L problem, not a taste complaint

You can already see the tell: a centered hero in Inter, an indigo-to-lavender gradient, three rounded cards, a badge above the headline. One researcher scored 1,590 real "Show HN" launches against 16 deterministic checks for this fingerprint and found 22% heavily slop-marked, 32% mildly so — a majority of real, funded products wearing the same uniform. It traces to a specific, traceable cause: Tailwind's creator picked `indigo-500` as a neutral demo accent in 2020, it saturated years of tutorials and starter templates, and today's models learned the correlation between "button" and that exact color. He publicly apologized for it in 2025. The mechanism is general — models return the highest-probability output for an under-specified prompt, and the highest-probability output is whatever the internet had the most of.

This matters commercially, not just aesthetically, for four concrete reasons: users form a lasting first impression of a page in about 50 milliseconds, before they've read a word, so an indistinguishable product has already lost the one moment it had to be remembered; an unedited AI default signals to a visitor that nobody at the company cared enough to override it; two of the most common "slop tells" — low-contrast dark-mode text and forced dark mode with no alternative — are outright accessibility failures, and the EU's Accessibility Act has been legally enforceable since June 2025; and when every competitor can generate the same competent-looking product for free, you're no longer competing on differentiation, you're competing on price. There's a real, measured cost to the *hollow* version of this too: a 2025 HBR/BetterUp study of 1,150 workers put the cost of "workslop" — AI output that looks like finished work but isn't — at roughly $186 per employee per month, north of $9M a year for a 10,000-person company.

Where AI actually breaks — and where it doesn't

It's worth being precise about the ceiling, because it tells you exactly what to build around it. The best current vision models judge "is this design good" correctly about 72% of the time, and asked to draw a box around exactly what's wrong, they're barely better than random — reasoning models don't help, and sometimes score *worse*. Left to iterate on their own without a corrective signal, generative loops reliably collapse: one controlled study ran 700 independent image-generation loops across seven different "creativity" settings and found all 700 converged into just twelve generic visual themes — turning up the temperature dial didn't stop it. The root cause isn't the model's raw training; it's the later alignment step, which measurably rewards familiar, predictable output over novel output.

The useful way to split this is four separable capabilities, only two of which are actually scarce: **perception** (can a system tell slop when it sees it?) is close to commodity today — off-the-shelf linters already catch most of it deterministically. **Generation** (can it produce a competent screen?) is fully commoditized — that's what v0, Lovable, and Bolt sell for free or near-free. **Judgment** (can it decide a competent option is wrong for *this* brand, and refuse it?) is genuinely scarce, because it doesn't fall out of a prompt. And **coherence at scale** — holding a specific, non-default design language consistently across the 50th screen a growing team generates, not just the demo's first one — is scarce for the same reason: it requires a persistent memory of decisions that a chat-window context can't hold. The winning architecture, confirmed across every craft area in this research — color, type, layout, copy — is not "a smarter model." It's cheap deterministic rules for everything that's actually checkable (contrast ratios, spacing grids, color-token conformance), a small trained scoring model for "does this look like our quality bar," and a human for the residual judgment calls. That combination is buildable today; "an AI with taste" isn't.

The competitive map: where the money and the risk actually sit

By 2026 the category has split into roughly six overlapping lanes, and they're all converging toward the same endpoint — prompt in, running product out — which means *category* alone is not a moat. Prompt-to-app builders (v0, Bolt, Lovable, Replit) are the loudest and most commoditized; even the biggest of them are quietly moving off flat subscription pricing toward metered usage, because inference cost is variable and app value is spiky. Prompt-to-*design* tools are being absorbed by incumbents outright: Miro bought Uizard, and Google turned the well-funded startup Galileo AI into a free Google Labs product (Stitch), resetting that entire category's price to zero — a hyperscaler can give away in a day what a startup was charging for. Figma, the incumbent with the most to defend, went public in mid-2025 at $274M quarterly revenue and is responding on every front at once: its own prompt-to-app tool (Figma Make), an on-canvas agent, an acquisition, and — tellingly — bolting *metered AI credits* onto its seat-based pricing, arriving at the same seats-plus-usage hybrid as the insurgents, just from the other direction. Design-to-code specialists (Anima, Locofy, Builder.io) all land around "roughly 75% of the way there" on real production handoff and sell mainly to enterprise teams, because that's where the value — a maintainable codebase inside a real design system — is durable enough to pay for.

An independent Nielsen Norman Group evaluation of a dozen-plus leading tools in late 2025 found the same failures across the board: poor contrast, inconsistent spacing, and "a similar, generic look" — quality only improved with long, detailed prompts, meaning most of the actual design decision-making had to happen *before* the AI touched it. That's the diagnostic for where to build: **single-screen generation is commodity and getting cheaper. Design-system fidelity, dense professional interfaces, and**

production-grade handoff are still hard, and that's where a real product lives. Value is shifting from *generating* a look from nothing toward *extracting and enforcing* a brand's actual existing system — which is a fundamentally different, more defensible product than a better generator.

Where the moat actually is

Start from an uncomfortable data point: Hugging Face reproduced OpenAI's flagship Deep Research agent to 82% of its benchmark score, open-source, in **24 hours**. Whatever your demo does today, a competitor's does in weeks. The canonical cautionary tale is Jasper, which raised $125M at a $1.5B valuation in October 2022 on GPT-3 plus copywriting templates — five weeks before ChatGPT launched and commoditized the entire premise overnight; its revenue fell from a ~$120M peak toward ~$55M. If your product's value is "a capability the model vendor is about to give everyone for free," you don't have a business, you have a head start.

Five things actually hold up as moats, and foundation-model access is explicitly not one of them, because everyone has the same APIs. **Workflow depth** — becoming the surface where the real work and real decisions happen, not a generator bolted onto the side of someone else's workflow. **Proprietary feedback data** — but with a real asterisk: most "data moats" are just scale effects that flatten out fast; it only defends when the data is genuinely proprietary to your product (not scraped from the open web everyone already trained on), the domain has real accuracy stakes, and it required earned trust to collect. **Distribution** — Microsoft's Copilot sits inside 90% of the Fortune 100, which is a real threat to any better-built competitor without that reach; concede this one rather than out-argue it. **Switching costs**, meaning genuinely accumulated context that makes leaving lossy — not lock-in for its own sake. And **governed data with real risk transfer** — Adobe trains Firefly only on licensed content and offers enterprise IP indemnification, a legal and provenance posture a scrappier competitor structurally can't match cheaply.

The practical version of "proprietary feedback data" is the one worth building toward deliberately: taste becomes a real asset the moment it stops being a founder's personal eye and becomes a dataset plus an evaluation system, and both compound with usage. Every accept, reject, and — most valuably — human edit to a generated design is a labeled preference signal. This doesn't require huge scale to start: one landmark study built a genuinely useful critique model from just seven professional designers' curated feedback; Midjourney's personalization stabilizes after roughly 200 pairwise ratings per user. Low thousands of curated judgments, not millions, is enough to get a first version working — and then usage compounds it. One governance problem has to be solved before your first enterprise deal, though: a customer's proprietary brand system must never leak into a shared model that also serves their competitors. The workable answer is two-tiered — a customer's own tokens and brand rules stay tenant-scoped and never cross-train; only the domain-general layer ("what kinds of things get rejected across our whole vertical") compounds across customers.

Before you build, answer four questions honestly. Is "good" definable in your target domain — checkable, not just felt? Can you own the actual feedback loop — the moment someone accepts, edits, or rejects, not just the final exported file? Can you enforce craft *deterministically,* so quality holds regardless of which underlying model generated the output, rather than depending on the model behaving well? And after six months, is there something real a customer loses by leaving — accumulated brand knowledge, not just a subscription? The first and third are hard gates: a "no" on either one means you don't have a moat, no matter how much usage data you eventually collect. The other two are fixable over time.

A short practical checklist

Before writing code: can you name the specific vertical and workflow you sit inside (not "an AI design tool")? Is quality in that domain at least partly checkable — contrast, brand-token conformance, layout rules — rather than pure unstructured taste? Can you assemble or license an exemplar corpus a competitor can't just scrape from the same public web every model already trained on?

In the product itself: never let the model invent raw color or spacing values — it should choose *from* a locked palette and scale, not improvise one. Run a slop-fingerprint check (Inter by default, the indigo/purple gradient, three identical cards, a badge over a centered headline) as an automated gate before anything ships. Treat contrast, spacing-grid adherence, and accessibility as hard, cheap, automatic checks — never something you ask a model to eyeball. Build the review/accept/reject moment into your own product as a first-class step, because that's where your proprietary training data actually gets created — if the file just gets exported and disappears, you never accumulate the one asset that compounds.

The closing test

Ask yourself one question honestly: if a competitor got your exact model access and your exact prompts tomorrow morning, what would they still be missing? A good answer is your accumulated exemplar and preference data, your deterministic quality guardrails, your specific customers' accumulated brand state, or your distribution. A bad answer is your prompts, your choice of model, or how polished your UI looks today — all three of those are commodity, and all three will be matched in weeks. The craft this research surfaced — in color, type, layout, and copy — was never decoration. Encoded as constraints, checks, and a feedback loop that gets smarter with every customer decision, it's the actual inventory of a defensible business. Everything else is renting the median.